

STATUS: DEPLOYED ●

Your agent is live. You have no idea if it's working.

Thread length is low. Retention is dropping. No error fired. No one complained loud enough. You just have no signal.



No errors doesn't mean it's working.

You grabbed a third-party eval tool. Relevance: 8/10. Coherence: 9/10. Looks like signal. But it knows nothing about your domain, your users, or what a correct answer even looks like.

RELEVANCE

8/10

GENERIC

COHERENCE

9/10

GENERIC

Some failures your code already sees.

Exceptions, tool call failures, malformed JSON structured outputs. Write fast boolean assertions on format, schema, and tool responses. Run on every change. No LLM needed.

 test_agent_outputs.py

```
assert is_valid_json(response)
```

PASS

```
assert tool_call.status == 200
```

PASS

```
assert schema.validate(output)
```

PASS

```
assert no_exception_raised()
```

FAIL


Your chat logs are a passive eval you never ran.

Classify frustrated conversations — short threads, repeated rephrasing, abrupt drop-offs. These users were already telling you what’s broken. Cluster the failure patterns that surface.



Pull top 200 queries. Let an expert label 50–100.

Not you the engineer — a domain expert who knows what a correct answer looks like. Binary labels, one-line reasoning each. That dataset is your golden standard for everything that follows.



Now build the judge that thinks like your expert.

Run a judge prompt on the labeled pairs. Measure agreement score. Iterate until it matches expert judgment. Fix the agent. Run the eval — numbers move. Unit test the fix so it never regresses.

AGREEMENT // BEFORE

61%

judge $\neq$ expert



AGREEMENT // AFTER

94%

locked w/ unit test

How did your eval setup evolve in the last AI agent you worked with?

Drop it in the comments — zero evals, fancy dashboard, or something in between.

☐ zero evals

☐ fancy dashboard

☒ something in between

> 